

## Question 1

### 1. Output of a single head.

For head  $k \in \{1, \dots, K\}$ , we first compute unnormalized attention scores

$$e_{ij}^{(t+1,k)} = \text{LeakyReLU}\left(\left(\mathbf{a}^{(k)}\right)^\top [\mathbf{W}^{(k)} \mathbf{z}_i^{(t)} \| \mathbf{W}^{(k)} \mathbf{z}_j^{(t)}]\right),$$

and normalize them over the neighbors of  $i$ :

$$\alpha_{ij}^{(t+1,k)} = \frac{\exp(e_{ij}^{(t+1,k)})}{\sum_{p \in \mathcal{N}(i)} \exp(e_{ip}^{(t+1,k)})}.$$

The output of head  $k$  for node  $i$  is then

$$\mathbf{z}_i^{(t+1,k)} = \sigma\left(\sum_{j \in \mathcal{N}(i)} \alpha_{ij}^{(t+1,k)} \mathbf{W}^{(k)} \mathbf{z}_j^{(t)}\right),$$

where  $\sigma(\cdot)$  is a nonlinearity (e.g. ELU or ReLU).

### 2. Concatenation over $K$ heads.

The final representation of node  $i$  is obtained by concatenating the outputs of all heads:

$$\mathbf{z}_i^{(t+1)} = \left\|_{k=1}^K \mathbf{z}_i^{(t+1,k)} \right\|.$$

Since each head produces a vector in  $\mathbb{R}^{F'_\text{out}}$ , the total dimensionality of  $\mathbf{z}_i^{(t+1)}$  is

$$\dim(\mathbf{z}_i^{(t+1)}) = K F'_\text{out}.$$

### 3. Total number of learnable parameters.

For a single head  $k$ :

- The weight matrix  $\mathbf{W}^{(k)} \in \mathbb{R}^{F'_\text{out} \times F_\text{in}}$  contributes  $F'_\text{out} F_\text{in}$  parameters.
- The attention vector  $\mathbf{a}^{(k)} \in \mathbb{R}^{2F'_\text{out}}$  contributes  $2F'_\text{out}$  parameters.

Hence, one head has

$$F'_\text{out} F_\text{in} + 2F'_\text{out} = F'_\text{out}(F_\text{in} + 2)$$

parameters. With  $K$  independent heads, the total number of parameters (ignoring biases for simplicity) is

$$K F'_\text{out}(F_\text{in} + 2).$$

## Question 2

We assume that all nodes share the same feature vector:

$$\mathbf{x}_i = \mathbf{c} \in \mathbb{R}^d \quad \text{for all } v_i \in V.$$

### 1. Unnormalized score $e_{ij}$ and coefficient $\alpha_{ij}$ .

In a single-head GAT layer, we have

$$e_{ij} = \text{LeakyReLU}\left(\mathbf{a}^\top [\mathbf{W} \mathbf{x}_i \| \mathbf{W} \mathbf{x}_j]\right), \quad \alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k \in \mathcal{N}(i)} \exp(e_{ik})}.$$

Under the assumption  $\mathbf{x}_i = \mathbf{c}$  for every  $i$ , we have

$$\mathbf{Wx}_i = \mathbf{Wc}, \quad \mathbf{Wx}_j = \mathbf{Wc},$$

which implies

$$e_{ij} = \text{LeakyReLU}\left(\mathbf{a}^\top [\mathbf{Wc} \parallel \mathbf{Wc}]\right) =: \tilde{e},$$

where  $\tilde{e}$  is a constant (independent of  $i$  and  $j$ ). Therefore, for any neighbor  $j \in \mathcal{N}(i)$ :

$$\alpha_{ij} = \frac{\exp(\tilde{e})}{\sum_{k \in \mathcal{N}(i)} \exp(\tilde{e})} = \frac{\exp(\tilde{e})}{|\mathcal{N}(i)| \exp(\tilde{e})} = \frac{1}{|\mathcal{N}(i)|}.$$

## 2. Effect on the expressiveness of the GAT layer.

All neighbors of a node receive the same attention weight:

$$\alpha_{ij} = \frac{1}{|\mathcal{N}(i)|}.$$

The update rule becomes

$$\mathbf{z}_i^{(t+1)} = \sigma\left(\frac{1}{|\mathcal{N}(i)|} \sum_{j \in \mathcal{N}(i)} \mathbf{Wz}_j^{(t)}\right),$$

i.e. a simple average of transformed neighbor features. The mechanism no longer distinguishes between neighbors, so it no longer behaves as a meaningful *attention* mechanism.

In this limit, the layer effectively reduces to a standard neighborhood averaging GNN, similar in spirit to a mean-aggregator GraphSAGE layer or a simplified GCN layer.

## 3. Why can it still beat random guessing on the Karate graph?

Even though the node features carry no useful information and attention becomes uniform, the model still operates on a non-trivial graph structure: the connectivity pattern in the Karate club network is highly correlated with the community labels.

## Question 3

To obtain a *conditional* VGAE, we introduce an additional conditioning variable  $\mathbf{c}$  (for instance a graph-level label, side information, or some global attributes). Instead of modeling

$$q_\phi(\mathbf{Z} | A, X) \quad \text{and} \quad p_\theta(A | \mathbf{Z}),$$

we make both the encoder and decoder explicitly depend on  $\mathbf{c}$ :

$$q_\phi(\mathbf{Z} | A, X, \mathbf{c}), \quad p_\theta(A | \mathbf{Z}, \mathbf{c}),$$

and possibly define a conditional prior  $p(\mathbf{Z} | \mathbf{c})$ .

From an architectural viewpoint, this can be implemented in several ways:

- **Encoder:** concatenate  $\mathbf{c}$  to node features or graph embeddings before the encoder GNN, or inject it via conditioning mechanisms (e.g. FiLM layers). Symbolically,

$$X' = [X \parallel \mathbf{c}], \quad q_\phi(\mathbf{Z} | A, X, \mathbf{c}) = q_\phi(\mathbf{Z} | A, X').$$

- **Decoder:** condition the adjacency reconstruction on both  $\mathbf{Z}$  and  $\mathbf{c}$ , for example by concatenating  $\mathbf{c}$  to graph-level latent codes before the decoder MLP, or using  $\mathbf{c}$  to modulate decoder layers:

$$A \sim p_\theta(A | \mathbf{Z}, \mathbf{c}).$$

- **Prior:** instead of a standard isotropic prior, use

$$p(\mathbf{Z} | \mathbf{c}) = \mathcal{N}(\mu(\mathbf{c}), \Sigma(\mathbf{c})),$$

where  $\mu(\mathbf{c})$  and  $\Sigma(\mathbf{c})$  are learned functions of the condition.

## Question 4

In a VAE/VGAE, the encoder outputs the parameters  $(\mu, \sigma)$  of a Gaussian distribution and we want to sample

$$\mathbf{z} \sim \mathcal{N}(\mu, \sigma^2 I)$$

while still being able to backpropagate through  $\mu$  and  $\sigma$ .

### Reparameterization trick

We rewrite the sampling step as

$$\mathbf{z} = \mu + \sigma \odot \epsilon, \quad \epsilon \sim \mathcal{N}(0, I),$$

so that  $\epsilon$  is independent noise and  $\mathbf{z}$  is a differentiable function of  $(\mu, \sigma)$  for a fixed draw of  $\epsilon$ . Standard backpropagation can then be used to obtain gradients with respect to  $\mu$  and  $\sigma$ .

### Without reparameterization

If we instead sampled

$$\mathbf{z} \sim \mathcal{N}(\mu, \sigma^2 I)$$

directly via a black-box sampler,  $\mathbf{z}$  would not be differentiable with respect to  $\mu$  and  $\sigma$ . No useful gradients could flow from the loss to the encoder, and one would need high-variance estimators. The reparameterization trick avoids this by making the sampling step itself differentiable, allowing stable end-to-end training by backpropagation.